

# Ensemble

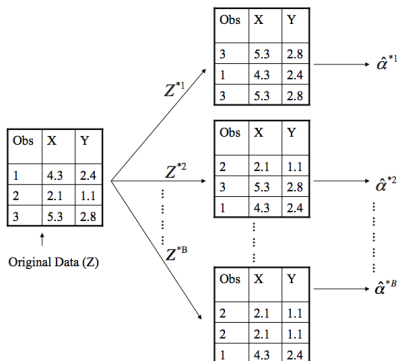
Yang Wang

College of Charleston

- Decision trees are simple and interpretable models for regression and classification.
- However, they are often not competitive with other methods in terms of prediction accuracy
- Bagging (Bootstrap aggregating), random forest, and boosting are good methods for improving the prediction accuracy of trees. They work by growing many trees on the training data and then combining the predictions of the resulting ensemble of trees.
- The latter two methods - random forests and boosting - are among the state-of-the-art methods for supervised learning, However, their results can be difficult to interpret.

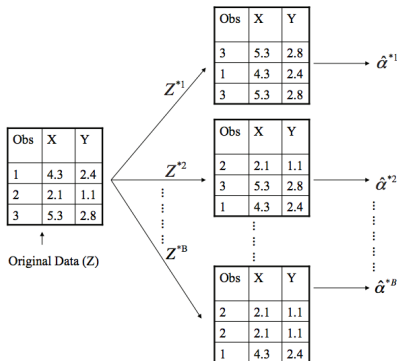
# Bootstrapping

- Resampling procedure
- Idea: repeatedly draw samples with replacement from a known sample and compute the item of interest on each sample



# Bootstrapping: Example

- Original dataset has 3 observations
- $B$  bootstrap samples (each of size 3) are created by drawing sample with replacement from the original dataset
- Suppose  $\alpha$  is a parameter we are interested in (example: mean)
- We will compute an estimate of  $\alpha$  *i.e.*  $\hat{\alpha}$  for each bootstrap sample
- The  $B$  estimates of  $\alpha$  obtained from  $B$  bootstrap samples gives a sampling distribution of  $\hat{\alpha}$



# Bootstrapping: Example


- Suppose we have the following data on the height of 15 individuals
- Data (in inches): 72 70 65 71 72 66 60 61 66 62 67 67 68 67 71
- We want to create a confidence interval for the mean height
- Draw 2000 samples of size 15 with replacement - these are the bootstrap samples

```
> ## original data
> t(height)
[,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10] [,11] [,12] [,13] [,14] [,15]
63   64   67   71   62   71   71   68   68   61   62   62   68   65   69
> ## create 2000 bootstrap samples
> bt_samples <- bootstraps(height, times = 2000)
> ## example of first bootstrap sample
> height[bt_samples$splits[1][[1]]$in_id,]
[1] 68 62 69 71 62 69 71 61 64 62 71 63 71 65 71
> ## example of second bootstrap sample
> height[bt_samples$splits[2][[1]]$in_id,]
[1] 68 68 68 67 68 62 62 64 62 71 68 61 62 68 68
```

# Bootstrapping: Example

- Compute mean on each bootstrap sample.
- The histogram shows the sampling distribution of mean.
- Find out the 0.025 and 0.975 sample quantiles of the bootstrap means.
- These sample quantiles constructs the 95% confidence interval for mean height.


```
## calculate mean on each bootstrap sample
bt_mean <- rep(NA, 2000)
for (i in 1:2000){
  bt_mean[i] <- mean(height[bt_samples$plits[i][[1]]$in_id,])
}
hist(bt_mean, main = "sampling distribution of bootstrap mean")
## 95% confidence interval for mean height using bootstrap method
> quantile(bt_mean, probs = c(0.025, 0.975))
 2.5%      97.5%
64.40000 67.93333
> ## traditional confidence intervals
> c(mean(height[,1]) - qt(.975,14)*sd(height[,1])/sqrt(15),
    mean(height[,1]) + qt(.975,14)*sd(height[,1])/sqrt(15))
[1] 64.12716 68.13951
```



- Advantage: no distribution is assumed

# Ensembles

- Goal: predict  $Y$  for given  $X$
- May be used for both classification and regression
- Combines the results of multiple models

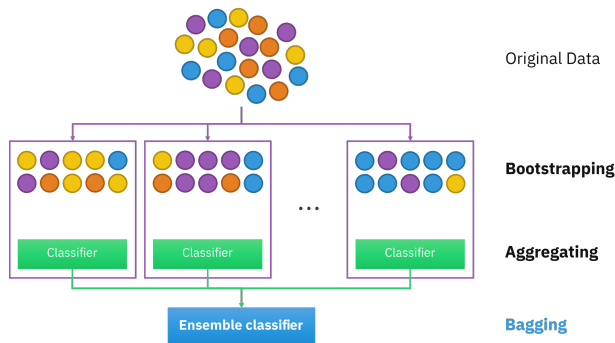


picture:

<https://www.analyticsvidhya.com/blog/2018/06/comprehensive-guide-for-ensemble-models/>

# Bagging (Bootstrap aggregating)

- Generate  $B$  bootstrap samples.
- Fit predictive model on each bootstrap sample.
- The final prediction is an average of the predictions obtained from the  $B$  models fitted on bootstrap samples.
  - Regression:  $\frac{1}{B} \sum_{b=1}^B T_b(\mathbf{x})$
  - Classification: choose the majority class label among  $T_1(\mathbf{x}), T_2(\mathbf{x}), \dots, T_B(\mathbf{x})$ .





# Reduce Variance Without Increasing Bias

- Recall:  $E(X+Y) = E(X) + E(Y)$
- Recall:  $\text{Var}(X+Y) = \text{Var}(X) + \text{Var}(Y) + 2 \text{Cov}(X,Y)$
- Recall:  $\text{Var}(aX) = a^2 \text{Var}(X)$  ( $a$  is a constant)
- Consider identically distributed random variables  $T_1, T_2, \dots, T_B$  which are used for estimating a parameter  $\theta$ .
- The bias of the statistic  $T_i$  is given by  $\text{bias}(\theta) = E(T_i) - \theta$ . If  $\text{bias}(\theta) = 0$ .
- $T_i$  is called an unbiased estimator of parameter  $\theta$ .
- Define  $T_{\text{avg}} = \frac{1}{B}(T_1 + \dots + T_B)$ .
- $E(T_{\text{avg}}) = \frac{1}{B} \sum_{i=1}^B E(T_i)$ .
- $\text{Var}(T_{\text{avg}}) = \frac{1}{B^2} (\sum_{i=1}^B \text{Var}(T_i) + \sum_{i=1}^B \sum_{\substack{j=1 \\ i \neq j}}^B \text{Cov}(T_i, T_j))$ .

# Reduce Variance Without Increasing Bias

- $E(T_{avg}) = \frac{1}{B} \sum_{i=1}^B E(T_i).$
- $Var(T_{avg}) = \frac{1}{B^2} (\sum_{i=1}^B Var(T_i) + \sum_{i=1}^B \sum_{\substack{j=1 \\ i \neq j}}^B Cov(T_i, T_j)).$
- Suppose  $E(T_i) = \mu$ ,  $Var(T_i) = \sigma^2$  and  $Cov(T_i, T_j) = \rho\sigma^2$
- $\rho$  is the correlation between  $T_i$  and  $T_j$ .
- $E(T_{avg}) = \mu$  i.e bias remains the same.
- $Var(T_{avg}) = \frac{1}{B^2} (B\sigma^2 + B(B-1)\rho\sigma^2) = \rho\sigma^2 + (1-\rho)\sigma^2/B.$
- For sufficiently large  $B$  the second term vanishes.
- Since  $\rho < 1$ , the variance of  $Var(T_{avg}) < \sigma^2$ .
- Averaging reduces variance while retaining the bias.
- Idea of ensemble method: build multiple models and use their average for variance reduction.

# Bagging: R example

## Load the diabetes.csv

```
library(caret)
set.seed(0)
train_control <- trainControl(method="cv", number=10)
## specify nbagg to control the number of trees. default value is 25
bag <- train(Outcome ~ ., data = training.data, method = "treebag",
            trControl = train_control, nbagg = 50)

> print(bag)
Bagged CART

538 samples
   8 predictor
   2 classes: 'no', 'yes'

No pre-processing
Resampling: Cross-Validated (10 fold)
Summary of sample sizes: 484, 484, 484, 484, 484, 485, ...
Resampling results:

   Accuracy   Kappa
0.7510482  0.433193

plot(varImp(bag))
> bag$finalModel

Bagging classification trees with 50 bootstrap replications
```

# Bagging: R example

```
# get prediction on the test data
pred.test.bag = predict(bag$finalModel, test.data, type = 'class')

# create confusion matrix
> confusionMatrix(pred.test.bag, test.data$Outcome, positive = "yes")
```

Confusion Matrix and Statistics

	Reference	
Prediction	no	yes
no	121	26
yes	29	54

Accuracy : 0.7609  
95% CI : (0.7004, 0.8145)  
No Information Rate : 0.6522  
P-Value [Acc > NIR] : 0.000245

Kappa : 0.4775


McNemar's Test P-Value : 0.787406

Sensitivity : 0.6750  
Specificity : 0.8067  
Pos Pred Value : 0.6506  
Neg Pred Value : 0.8231  
Prevalence : 0.3478  
Detection Rate : 0.2348  
Detection Prevalence : 0.3609  
Balanced Accuracy : 0.7408

'Positive' Class : yes

# Bagging: R example

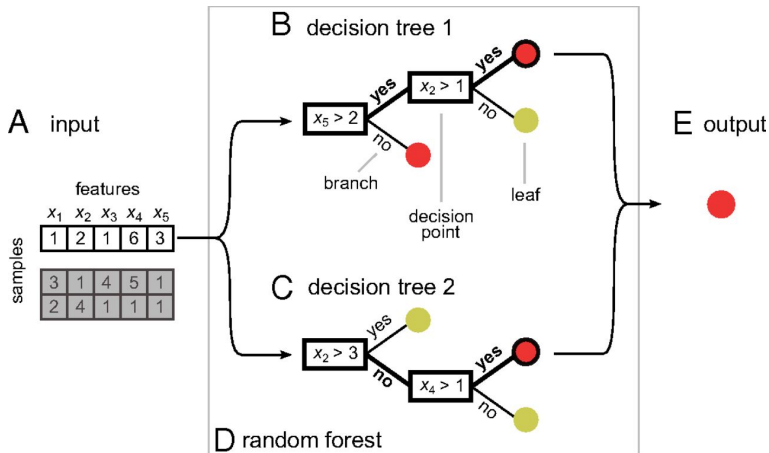
## Variable importance plot



# Random forest

- Generate  $B$  bootstrap samples.
- Fit decision trees on each bootstrap sample - but with a little tweak.
  - Each time a split in a tree is considered, do not use the full set of  $p$  predictors.
  - Use a random sample of  $m$  predictors (typically  $m \approx \sqrt{p}$ ).
- The final prediction is an average of the predictions obtained from the  $B$  models fitted on bootstrap samples.
  - Regression:  $\frac{1}{B} \sum_{b=1}^B T_b(\mathbf{x})$
  - Classification: choose the majority class label among  $T_1(\mathbf{x}), T_2(\mathbf{x}), \dots, T_B(\mathbf{x})$ .

# Random forest



picture: <https://www.pnas.org/content/115/8/1690>

# Why does random forest work?

- Random forest often performs better than bagging, although at each split in the tree, the algorithm is not even allowed to consider a majority of the available predictors.
- Suppose that there is one very strong predictor in the data set, along with a number of other moderately strong predictors.
- Then in the collection of bagged trees, most or all of the trees will use this strong predictor in the top split.
- All of the bagged trees will look quite similar to each other and the predictions from the bagged trees will be highly correlated.
- Averaging many highly correlated quantities does not lead to as large of a reduction in variance as averaging many uncorrelated quantities.
- This means that bagging will not lead to a substantial reduction in variance over a single tree in this setting.
- Random forests overcome this problem by forcing each split to consider only a subset of the predictors.
- Therefore, on average  $(p - m)/p$  of the splits will not even consider the strong predictor, and so other predictors will have more of a chance.
- We can think of this process as decorrelating the trees, thereby making the average of the resulting trees less variable and hence more reliable.



# Random forest: R example

```
library(caret)
set.seed(0)
> modelLookup("rf")
  model parameter                                label forReg forClass probModel
1    rf          mtry #Randomly Selected Predictors      TRUE      TRUE      TRUE

train_control <- trainControl(method="cv", number=10)
rf <- train(Outcome ~ . , data = training.data, method = "rf",
           trControl = train_control, tuneLength = 3)
# metric = "Kappa" may be mentioned if best tree should be selected based on that
## specify tuning parameters using tuneGrid
# rf <- train(Outcome ~ . , data = training.data, method = "rf", ntree = 50,
#           trControl = train_control, tuneGrid = expand.grid(mtry = c(2,5,8)))
> print(rf)
```

Random Forest

```
538 samples
  8 predictor
  2 classes: 'no', 'yes'
```

No pre-processing

Resampling: Cross-Validated (10 fold)

Summary of sample sizes: 484, 484, 484, 484, 484, 485, ...

Resampling results across tuning parameters:

mtry	Accuracy	Kappa
2	0.7492313	0.4221906
5	0.7603774	0.4519239
8	0.7492662	0.4330171

Accuracy was used to select the optimal model using the largest value.

The final value used for the model was mtry = 5.

# Random forest: R example

```
> rf$finalModel
Call:
  randomForest(x = x, y = y, mtry = param$mtry)
      Type of random forest: classification
      Number of trees: 500
No. of variables tried at each split: 5
      OOB estimate of  error rate: 25.84%
Confusion matrix:
      no yes class.error
no  292  58  0.1657143
yes  81 107  0.4308511

# get prediction on the test data
pred.test.rf = predict(rf$finalModel, test.data, type = 'class')
# create confusion matrix
> confusionMatrix(pred.test.rf, test.data$Outcome, positive = "yes")
Confusion Matrix and Statistics

      Reference
Prediction no yes
      no  125  23
      yes  25  57

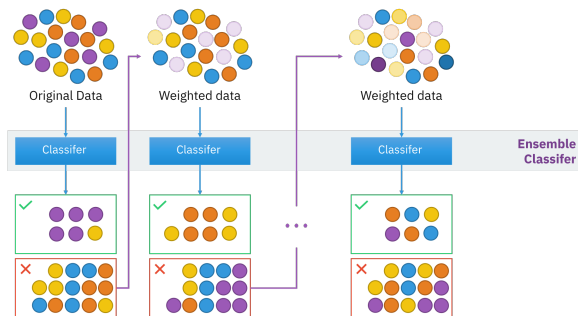
      Accuracy : 0.7913
      95% CI   : (0.733, 0.8419)
      No Information Rate : 0.6522
      P-Value [Acc > NIR] : 2.878e-06
      Kappa   : 0.5427
      McNemar's Test P-Value : 0.8852
      Sensitivity : 0.7125
      Specificity : 0.8333
      Pos Pred Value : 0.6951
      Neg Pred Value : 0.8446
      Prevalence : 0.3478
```

# Boosting

- Recall: Bagging
  - Generate  $B$  bootstrap samples
  - Fit model on each bootstrap sample
  - Combine all models to create a simple predictive model
  - Each tree is built on a bootstrap data set, which is independent of the other trees
- Boosting works in a similar way, except that the trees are grown sequentially
- Each tree is grown using information from previously grown trees
- Boosting does not involve bootstrap sampling; instead each tree is fit on a modified version of the original data set


# Boosting

- In the beginning the process assigns equal weight to each observations
- After fitting a tree, the performance of the model is evaluated
- The missclassified observations are assigned a higher weight and the correctly classified observations get a lower weight
- A tree is fitted again on the reweighted data
- The process is repeated M times
- The final prediction is a weighted average of the predictions of all fitted models



# Adaboost

- Initialize the observation weights  $w_i = 1/N, i = 1, 2, \dots, N$
- For  $m = 1$  to  $M$  :
  - Fit a classifier  $G_m(x)$  to the training data using  $w_i$
  - Compute  $err_m = \frac{\sum_{i=1}^N w_i \mathbf{1}(y_i \neq G_m(x_i))}{\sum_{i=1}^N w_i}$
  - Compute  $\alpha_m = \log((1 - err_m)/err_m)$
  - Set  $w_i \leftarrow w_i \cdot \exp[\alpha_m \cdot \mathbf{1}(y_i \neq G_m(x_i))], i = 1, 2, \dots, N$
- Output  $G(x) = \text{sign}[\sum_{m=1}^M \alpha_m G_m(x)]$



# Adaboost: R example

```
library(caret)
library(ada)
set.seed(0)
train_control <- trainControl(method="cv", number=10)
ada <- train(Outcome ~ ., data = training.data, method = "ada",
            trControl = train_control, tuneLength = 3)
> print(ada)
Boosted Classification Trees
538 samples
  8 predictor
  2 classes: 'no', 'yes'
```

No pre-processing

Resampling: Cross-Validated (10 fold)

Summary of sample sizes: 484, 484, 484, 484, 484, 485, ...

Resampling results across tuning parameters:

maxdepth	iter	Accuracy	Kappa
1	50	0.7622292	0.4451545
1	100	0.7585255	0.4294762
1	150	0.7658980	0.4523454
2	50	0.7511530	0.4228118
2	100	0.7548567	0.4318264
2	150	0.7660377	0.4614452
3	50	0.7641160	0.4597316
3	100	0.7660377	0.4644768
3	150	0.7678546	0.4727207

Tuning parameter 'nu' was held constant at a value of 0.1

Accuracy was used to select the optimal model using the largest value.

The final values used for the model were iter = 150, maxdepth = 3 and nu = 0.1, the largest value.

The final value used for the model was mtry = 5.

# Adaboost: R example

```
# get prediction on the test data
pred.test.ada = predict(ada$finalModel, test.data)

# create confusion matrix
> confusionMatrix(pred.test.ada, test.data$Outcome, positive = "yes")
Confusion Matrix and Statistics
```

```
          Reference
Prediction no yes
no      122  29
yes     28  51
```

```
          Accuracy : 0.7522
          95% CI   : (0.6912, 0.8066)
No Information Rate : 0.6522
P-Value [Acc > NIR] : 0.0007075
```

```
          Kappa    : 0.4522
```

```
McNemar's Test P-Value : 1.0000000
```

```
          Sensitivity : 0.6375
          Specificity : 0.8133
          Pos Pred Value : 0.6456
          Neg Pred Value : 0.8079
          Prevalence    : 0.3478
          Detection Rate : 0.2217
          Detection Prevalence : 0.3435
          Balanced Accuracy : 0.7254
```

```
'Positive' Class : yes
```

# Advantage and disadvantage of ensemble methods

- Advantage:
  - Improved accuracy by combining together hundreds of models into a single prediction.
  - Low variance.
- Shortcoming:
  - Difficult to interpret: when we bag a large number of trees, it is not possible to represent the resulting prediction using a single tree.